

Importing the required libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
import tensorflow as tf
import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
```

Loading the dataset on colab

```
df = pd.read_csv('/content/drive/MyDrive/Crop Production data.csv')
```

Displaying the first 10 rows of the dataset

```
df.head(10)

{"type": "dataframe", "variable_name": "df"}
```

Info of the dataset

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 246091 entries, 0 to 246090
Data columns (total 7 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   State_Name      246091 non-null object  
 1   District_Name   246091 non-null object  
 2   Crop_Year       246091 non-null int64   
 3   Season         246091 non-null object  
 4   Crop            246091 non-null object  
 5   Area            246091 non-null float64  
 6   Production      242361 non-null float64  
dtypes: float64(2), int64(1), object(4)
memory usage: 13.1+ MB
```

Finding the mean of the production

```
mean_production = df['Production'].mean()

mean_production

582503.4422509808
```

Filling the null values with the mean of production

```
df['Production'] = df['Production'].fillna(value = mean_production)
```

Info after filling the null values

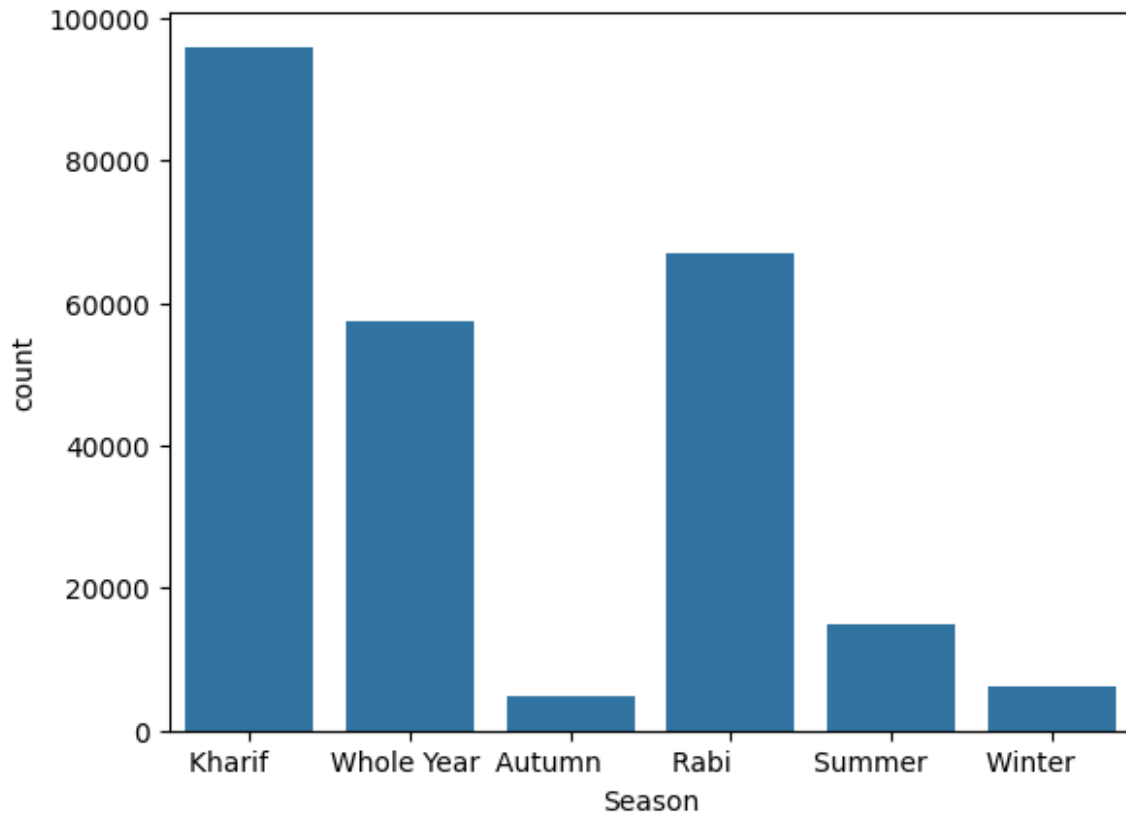
```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 246091 entries, 0 to 246090
Data columns (total 7 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   State_Name      246091 non-null object  
 1   District_Name   246091 non-null object  
 2   Crop_Year       246091 non-null int64   
 3   Season          246091 non-null object  
 4   Crop            246091 non-null object  
 5   Area            246091 non-null float64  
 6   Production      246091 non-null float64  
dtypes: float64(2), int64(1), object(4)
memory usage: 13.1+ MB
```

Countplot of the season


```
sns.countplot(x = 'Season', data = df)

<Axes: xlabel='Season', ylabel='count'>
```



values grouped by Season and displayed on the horizontal bar chart

```
from matplotlib import pyplot as plt
import seaborn as sns
df.groupby('Season').size().plot(kind='barh',
color=sns.palettes.mpl_palette('Dark2'))
plt.gca().spines[['top', 'right']].set_visible(False)
```




Plotting the series between Crop year and area

```
def _plot_series(series, series_name, series_index=0):
    palette = list(sns.palettes.mpl_palette('Dark2'))
    xs = series['Crop_Year']
    ys = series['Area']


    plt.plot(xs, ys, label=series_name, color=palette[series_index %
len(palette)])

fig, ax = plt.subplots(figsize=(10, 5.2), layout='constrained')
df_sorted = df.sort_values('Crop_Year', ascending=True)
for i, (series_name, series) in enumerate(df_sorted.groupby('Season')):
    _plot_series(series, series_name, i)
    fig.legend(title='Season', bbox_to_anchor=(1, 1), loc='upper left')
sns.despine(fig=fig, ax=ax)
plt.xlabel('Crop_Year')
_ = plt.ylabel('Area')
```



Plot of area

```
df['Area'].plot(kind='line', figsize=(8, 4), title='Area')
plt.gca().spines[['top', 'right']].set_visible(False)
```



```
df.head(15)
{"type": "dataframe", "variable_name": "df"}
```

Sum of production

```
print("Total Production:", df.Production.sum())
```

Total Production: 143348854606.9862

Total sum of area available

```
print("Total Area:", df.Area.sum())
```

Total Area: 2953786189.17

```
df.shape
```

(246091, 7)

Count of districts

```
df['District_Name'].value_counts()
```

District_Name	
BIJAPUR	945
TUMKUR	936
BELGAUM	925
HASSAN	895
BELLARY	887

...

HYDERABAD	8
KHUNTI	6
RAMGARH	6
NAMSAI	1
MUMBAI	1

Name: count, Length: 646, dtype: int64

Value counts with respect to how many times it is mention in the table according to the seasons in which crop production is done

```
df['State_Name'].value_counts()
```

State_Name	
Uttar Pradesh	33306
Madhya Pradesh	22943

Karnataka	21122
Bihar	18885
Assam	14628
Odisha	13575
Tamil Nadu	13547
Maharashtra	12628
Rajasthan	12514
Chhattisgarh	10709
Andhra Pradesh	9628
West Bengal	9613
Gujarat	8436
Haryana	5875
Telangana	5649
Uttarakhand	4896
Kerala	4261
Nagaland	3906
Punjab	3173
Meghalaya	2867
Arunachal Pradesh	2546
Himachal Pradesh	2494
Jammu and Kashmir	1634
Tripura	1412
Manipur	1267
Jharkhand	1266
Mizoram	957
Puducherry	876
Sikkim	714
Dadra and Nagar Haveli	263
Goa	208
Andaman and Nicobar Islands	203
Chandigarh	90

Name: count, dtype: int64

#Count of values per season

```
df['Season'].value_counts()
```

Season	
Kharif	95951
Rabi	66987
Whole Year	57305
Summer	14841
Winter	6058
Autumn	4949

Name: count, dtype: int64

#Finding the dummy values for the categorical variables

```
new_State = pd.get_dummies(df['State_Name'], dtype = 'int')
new_District = pd.get_dummies(df['District_Name'], dtype = 'int')
new_Season = pd.get_dummies(df['Season'], dtype = 'int')
new_Crop = pd.get_dummies(df['Crop'], dtype = 'int')
```

Converting the dummy values into the dataframe

```
State_ = pd.DataFrame(new_State)
District_ = pd.DataFrame(new_District)
Season_ = pd.DataFrame(new_Season)
Crop_ = pd.DataFrame(new_Crop)
```

#Concatenating the new dataframes with the original one

```
new_df = pd.concat([State_, District_, Season_, Crop_, df], axis = 1)
```

#Displaying the new dataframe

```
new_df.head(10)
{"type": "dataframe", "variable_name": "new_df"}
```

#Dropping the unnecessary attributes

```
new_df = new_df.drop({'State_Name', 'District_Name', 'Season',
'Crop'}, axis = 1)
new_df.head(10)
{"type": "dataframe", "variable_name": "new_df"}
```

#Selecting the predictor and the target variable

```
X = new_df.drop('Production', axis = 1)
y = new_df[['Production']]
```

#Printing the shape of predictor and target variable

```
print(X.shape)
print(y.shape)

(246091, 811)
(246091, 1)
```

#Scaling the predictor and target variable


```
Scaler = MinMaxScaler()
X = Scaler.fit_transform(X)
y = Scaler.fit_transform(y)
```

#Splitting the dataset into training and testing set

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.2)
```

#Printing the training and testing set

```
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(196872, 811)
(49219, 811)
(196872, 1)
(49219, 1)
```

#Developing the sequential model

```
model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Dense(units = 128, activation = 'relu',
input_shape = (811,)))
model.add(tf.keras.layers.Dense(units = 64, activation = 'relu'))
model.add(tf.keras.layers.Dense(units = 64, activation = 'relu'))
model.add(tf.keras.layers.Dropout(0.5))
model.add(tf.keras.layers.Dense(units = 32, activation = 'relu'))
model.add(tf.keras.layers.Dense(units = 32, activation = 'relu'))
model.add(tf.keras.layers.Dropout(0.5))
model.add(tf.keras.layers.Dense(units = 16, activation = 'relu'))
model.add(tf.keras.layers.Dense(units = 16, activation = 'relu'))
model.add(tf.keras.layers.Dropout(0.5))
model.add(tf.keras.layers.Dense(units = 8, activation = 'relu'))
model.add(tf.keras.layers.Dropout(0.5))
model.add(tf.keras.layers.Dense(units = 4, activation = 'relu'))
model.add(tf.keras.layers.Dense(units = 1, activation = 'linear'))
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	103936
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 64)	4160

dropout (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 32)	2080
dense_4 (Dense)	(None, 32)	1056
dropout_1 (Dropout)	(None, 32)	0
dense_5 (Dense)	(None, 16)	528
dense_6 (Dense)	(None, 16)	272
dropout_2 (Dropout)	(None, 16)	0
dense_7 (Dense)	(None, 8)	136
dropout_3 (Dropout)	(None, 8)	0
dense_8 (Dense)	(None, 4)	36
dense_9 (Dense)	(None, 1)	5

```
=====
Total params: 120465 (470.57 KB)
Trainable params: 120465 (470.57 KB)
Non-trainable params: 0 (0.00 Byte)
```

```
keras.utils.plot_model(model, to_file='png', show_shapes=True)
```

dense_input	input:	[(None, 811)]
InputLayer	output:	[(None, 811)]



dense	input:	(None, 811)
Dense	output:	(None, 128)



dense_1	input:	(None, 128)
Dense	output:	(None, 64)



dense_2	input:	(None, 64)
Dense	output:	(None, 64)



dropout	input:	(None, 64)
Dropout	output:	(None, 64)



dense_3	input:	(None, 64)
Dense	output:	(None, 32)



#Training the model with the help of callbacks so as to avoid overfitting

```
model.compile(optimizer = 'Adam', loss = 'mean_squared_error')
from keras.callbacks import EarlyStopping
es = EarlyStopping(patience = 2, monitor = 'val_loss')
model.fit(X_train, y_train, epochs = 25, batch_size = 10,
validation_data = (X_test, y_test), callbacks = [es])


Epoch 1/25
19688/19688 [=====] - 114s 6ms/step - loss:
1.6057e-04 - val_loss: 2.8405e-04
Epoch 2/25
19688/19688 [=====] - 107s 5ms/step - loss:
1.5994e-04 - val_loss: 2.8420e-04
Epoch 3/25
19688/19688 [=====] - 109s 6ms/step - loss:
1.6003e-04 - val_loss: 2.8405e-04

<keras.src.callbacks.History at 0x7b8566507370>
```

#Plotting the history of the model

```
hist = model.history.history
h = pd.DataFrame(hist)
h.plot()

<Axes: >
```




#Model predicting the values

```
y_predict = model.predict(X_test)
y_predict
1539/1539 [=====] - 3s 2ms/step
array([[0.00016082],
       [0.00016082],
       [0.00016082],
       ...,
       [0.00016082],
       [0.00016082],
       [0.00016082]], dtype=float32)
```


#Plot between predicted values and test values

```
plt.plot(y_test,y_predict, '^', color = 'r')
plt.xlabel('y_test')
plt.ylabel('y_predict')
Text(0, 0.5, 'y_predict')
```



#Plot between model's predictions and the true values

```
y_predict_original = Scaler.inverse_transform(y_predict)
y_test_original = Scaler.inverse_transform(y_test)
plt.plot(y_test_original, y_predict_original, '^', color = 'b')
plt.xlabel('model_predictions')
plt.ylabel('true_values')
Text(0, 0.5, 'true_values')
```



#Calculation of value of n

```
k = X_test.shape
k
n = len(X_test)
n
49219
```

#Calculation of Root Mean Squared Value

```
from sklearn.metrics import
r2_score,mean_squared_error,mean_absolute_error
from math import sqrt
RMSE =
float(format(np.sqrt(mean_squared_error(y_test_original,y_predict_orig
inal)), '0.3f'))
print(RMSE)
21080981.274
```

#Calculation of Mean Squared Value

```
MSE = mean_squared_error(y_test_original,y_predict_original)
print(MSE)
444407771476213.7
```

#Calculation of Mean Absolute Value

```
MAE = mean_absolute_error(y_test_original,y_predict_original)
print(MAE)
900126.1386866091
```

#Calculation of R2 Score

```
r2 = r2_score(y_test_original,y_predict_original)
print(r2)
-0.0006611469950694726
```